


getindata

HCatalog

Motivation



Inside a Data-Driven Company

- **Analysts use multiple tools for processing data**
 - Java MapReduce, Hive and Pig and more
- **Analysts create many (derivative) datasets**
 - Different formats e.g. CSV, JSON, Avro, ORC
 - Files in HDFS
 - Tables in Hive

We simply **pick a right tool and format** for each application!



Pig Analysts

- **To process data, they need to remember**
 - where a dataset is located
 - what format it has
 - what the schema is



```
songs = LOAD 'log' USING PigStorage(',') AS (artistName, songId, timestamp, user);
artists = FOREACH songs GENERATE artistName;
grouped = GROUP artists BY artistName;
counted = FOREACH grouped GENERATE group AS aritstName, COUNT(artists) AS cnt;
STORE counted INTO 'artist-count-pig';
]
```

Hive Analysts

- **They store popular datasets in (external) Hive tables**
- **To analyze datasets generated by Pig analysts, they need to know**
 - where a dataset is located
 - what format it has
 - what the schema is
 - how to load it into Hive table/partition

Changes, Changes And Changes

Let's start using more efficient format since today!



NO WAY!

We would have to **re-write, re-test and re-deploy** our applications!
This means a lot of **engineering work** for us!



MR, Pig, Hive And Data Storage

- **Hive reads data location, format and schema from metadata**
 - Managed by Hive Metastore
- **MapReduce encodes them in the application code**
- **Pig specifies them in the script**
 - Schema can be provided by the Loader

Conclusion

- **MapReduce and Pig are sensitive to metadata changes!**

Data Availability

Jeff: Is your dataset already generated?

Tom: Check in HDFS!

Jeff: What is the location?

Tom: /datastreams/20140101

Jeff: hdfs dfs -ls /datastreams/20140101

Not yet.... :(

Tom: Check it later!



HCatalog

- **Aims to solve these problems!**
- **First of all**
 - It knows the location, format and schema of our datasets!

HCatalog With Pig

■ Traditional way

```
raw = load '/datastreams/20140101' using MyLoader()  
      as (time:chararray, host:chararray, userId:int, songId:int, duration:int);  
...  
store output into '/data/users/20140101' using MyStorer();
```

HCatalog With Pig

■ Traditional way


```
raw = load '/datastreams/20140101' using MyLoader()  
    as (time:chararray, host:chararray, userId:int, songId:int, duration:int);  
...  
store output into '/data/users/20140101' using MyStorer();
```

■ HCatalog way

```
raw = load 'streams' using HCatLoader();  
...  
store output into users using HCatStorer('date=20140101');
```

Demo

Interacting with HCatalog



Demo

- 1. Upload a dataset in HDFS**
- 2. Add the dataset HCatalog**
- 3. Access the dataset with Hive and Pig**

HCatalog CLI

1. Upload a dataset in HDFS

```
$ hdfs dfs -put streams /data
```

Streams: timestamp, host, userId, songId, duration

2013-01-04 06:55:23	wa.stream.rock.net	224	4071	30
2013-01-07 06:45:32	ny.stream.rock.net	680	83	172
2013-01-08 18:22:22	phi.stream.rock.net	680	5534	72
2013-01-10 14:46:09	ny.stream.rock.net	30	4042	339
2013-01-10 23:08:34	ny.stream.rock.net	680	6284	306
2013-01-11 08:52:32	wa.stream.rock.net	869	2910	146

HCatalog CLI

2. Add the dataset HCatalog

- A file with HCatalog DDL can be prepared

```
CREATE EXTERNAL TABLE streams(time string, host string, userId int, songId int, duration int)
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE
LOCATION '/datastreams';
```

- And executed by hcat -f

```
$ hcat -f streams.hcatalog
```

HCatalog CLI

2. Add the dataset HCatalog

- Describe the dataset

```
$ hcat -e "describe streams"

OK

time          string          None
host          string          None
userid        int            None
songid        int            None
duration      int            None
```

HCatalog CLI

3. Access the dataset with Pig

```
raw_streams = LOAD 'streams' USING org.apache.hcatalog.pig.HCatLoader();  
all_count = FOREACH (GROUP raw_streams ALL) GENERATE COUNT(raw_streams);  
DUMP all_count;
```

```
$ pig -useHCatalog streams.pig  
...  
(93294)
```

HCatalog CLI


3. Access the dataset with Hive

```
$ hive -e "select count(*) from streams"
```

```
OK
```

```
93294
```

Goals




Three Main Goals

- 1. Provide an abstraction on top of datasets stored in HDFS**
 - Just use the name of the dataset, not the path
- 2. Enable data discovery**
 - Store all datasets (and their properties) in HCatalog
 - Integrate with Web UI
- 3. Provide notifications for data availability**
 - Process new data immediately when it appears

Supported Formats and Projects



Supported Projects And Formats



Custom Formats

- **A custom format can be supported**
 - But InputFormat, OutputFormat, and Hive SerDe must be provided

Pig Interface - HCatLoader

- **Consists of HCatLoader and HCatStorer**
- **HCatLoader read data from a dataset**
 - Indicate which partitions to scan by following the load statement with a partition filter statement

```
raw = load 'streams' using HCatLoader();  
valid = filter raw by date = '20140101' and isValid(duration);
```

Pig Interface - HCatStorer

- **HCatStorer writes data to a dataset**
 - A specification of partition keys can be also provided
 - Possible to write to a single partition or multiple partitions

```
store valid into 'streams_valid' using HCatStorer  
('date=20110924');
```

MapReduce Interface

- **Consists of HCatInputFormat and HCatOutputFormat**
- **HCatInputFormat accepts a dataset to read data from**
 - Optionally, indicate which partitions to scan
- **HCatOutputFormat accepts a dataset to write to**
 - Optionally, indicated with partition to write to
 - Possible to write to a single partition or multiple partitions

Hive Interface

- **There is no Hive-specific interface**
 - Hive can read information from HCatalog directly
- **Actually, HCatalog is now a submodule of Hive**


Conclusion

- **HCatalog enables non-Hive projects to access Hive tables**

Components

- **Hive Metastore to store information about datasets**
 - A table per dataset is created (the same as in Hive)
- **hcat CLI**
 - Create and drop tables, specify table parameters, etc
- **Programming interfaces**
 - For MapReduce and Pig
 - New ones can be implemented e.g. for Crunch
- **WebHCat server**
 - More about it later

Features



Data Discovery

- A nice web UI can be build on top of HCatalog
 - You have all Hive tables there for free!
 - See Yahoo!'s *illustrative* example below

The image contains two screenshots of the Yahoo! Global HCatalog Dashboard. The left screenshot, titled 'Discovery UI', shows a search bar with 'YAHOO!' and 'GLOBAL HCATALOG DASHBOARD'. It includes a dropdown menu 'The Best Cluster' showing 'Available Databases' such as 'audience_db', 'tumblr_db', 'user_db', 'adv_warehouse', and 'flickr_db'. Below this is a list of 'Available Tables (audience_db)' including 'page_clicks', 'ad_clicks', 'user_info', 'session_info', and 'audience_info'. The right screenshot, titled 'Table Display UI', shows a similar dashboard for the 'page_clicks' table. It displays 'Hcat Instance: The Best Cluster', 'Database: audience_db', 'Table: page_clicks', and 'Owner: Awesome Yahoo'. Below this is a table with columns 'Column', 'Type', and 'Description' for 'bcookie', 'timestamp', and 'uid'. The 'Description' column includes details like 'Standard browser cookie' and 'DD-MON-YYYY HH:MI:SS (AM/PM)'. At the bottom, it says '...more table information and properties (e.g. data format etc.)' and 'Partitions ...list of partitions'.

Pictures
come from
Yahoo's
presentation
at Hadoop
Summit San
Jose 2014

Properties Of Datasets

- **Can store data life-cycle management information**
 - Cleaning, archiving and replication tools can learn which datasets are eligible for their services

```
ALTER TABLE intermediate.featured-streams SET TBLPROPERTIES ('retention' = 30);

SHOW TBLPROPERTIES intermediate.featured-streams;
SHOW TBLPROPERTIES intermediate.featured-streams("retention");
```


Notifications

- **Provides notifications for certain events**
 - e.g. Oozie or custom Java code can wait for those events and immediately schedule tasks depending on them
- **Multiple events can trigger notifications**
 - A database, a table or a partition is added
 - A set of partitions is added
 - A database, a table or a partition is dropped

Evolution Of Data

- **Allows data producers to change how they write data**
 - No need to re-write existing data
 - HCatalog can read old and new data
 - Data consumers don't have to change their applications
- **Data location, format and schema can be changed**
 - In case of schema, new fields can be added

HCatalog Beyond HDFS




WebHCat Server

- Provides a REST-like web API for HCatalog
 - Send requests to get information about datasets

```
curl -s 'http://hostname:  
port/templeton/v1/ddl/database/db_name/table/table_name?user.  
name=adam'
```
 - Send requests to run Pig or Hive scripts
- Previously called *Templeton*

There Is More!



- **Data engineers, architects and instructors**
- **+4 years of experience in Apache Hadoop**
 - Working with hundreds-node Hadoop clusters
 - Developing Hadoop applications in many cool frameworks
 - Delivering Hadoop trainings for +2 years
- **Passionate about data**
 - Speaking at conferences and meetups
 - Blogging and reviewing books

