



Transactional Operations in Apache Hive

- Eugene Koifman

DataWorks Summit, San Jose 2018

Agenda

- A bit of history
- Current Functionality
- Design
- Future Plans
- Closing Remarks

Early Hive

- Transactions
 - ACID: Atomicity, Consistency, Isolation, Durability
- Atomicity - Rely on File System ‘rename’
 - Insert into T partition($p=1$) select - OK
 - Dynamic Partition Write – not OK
 - Multi-Insert statement – not OK 
 - FROM <expr> Insert into A select ... Insert Into B select ...
- Isolation - Lock Manager
 - S/X locks – not good for long running analytics 

Early Hive – Changing Existing Data

- Drop <...>
- Insert Overwrite = Truncate + Insert
 - Gets expensive if done often on small % of data

Goals

- Support ACID properties
- Support SQL Update/Delete/Merge
- Low rate of transactions
 - Not OLTP
 - Not a replacement for MySql or HBase

Features – Hive 3

Transactional Tables

- Not all tables support transactional semantics
 - Managed Tables
 - No External tables or Storage Handler (Hbase, Druid, etc)
- Fully ACID compliant
- Single statement transactions
- Cross partition/cross table transactions
- Snapshot Isolation
 - Between Serializable and Repeatable Read

Transactional Tables – Full CRUD

- ◆ Supports Update/Delete/Merge
- ◆ CREATE TABLE T(a int, b int) STORED AS ORC TBLPROPERTIES ('transactional'='true');
- Restrictions
 - Managed Table
 - Table cannot be sorted
 - Currently requires ORC File but anything implementing
 - AcidInputFormat/AcidOutputFormat
 - Bucketing is optional!
- If upgrading from Hive 2
 - Requires Major Compaction before Upgrading

Transactional Tables – Insert only

- ◆ CREATE TABLE T(a int, b int) TBLPROPERTIES ('transactional'='true', 'transactional_properties'='insert_only');
 - Managed Table
 - Any storage format

Transactional Tables – Convert from flat tables

- ◆ `ALTER TABLE T SET_TBLPROPERTIES ('transactional'='true')`
- ◆ `ALTER TABLE T(a int, b int) SET_TBLPROPERTIES ('transactional'='true', 'transactional_properties'='true');`
 - Metadata Only operation
 - Compaction will eventually rewrite the table

Transactional Tables - New In Hive 3

- Alter Table Add Partition...
- Alter Table T Concatenate
- Alter Table T Rename To....
- Export/Import Table
- Non-bucketed tables
- Load Data... Into Table ...
- Insert Overwrite
- Fully Vectorized
- Create Table As ...
- LLAP Cache
- Predicate Push Down

Design – Hive 3

Transactional Tables – Insert Only

- Transaction Manager
 - Begin transaction and obtain a Transaction ID
 - For each table, get a Write ID – determines location to write to

```
create table TM (a int, b int) TBLPROPERTIES  
('transactional'='true',  
'transactional_properties'='insert_only');
```

```
insert into TM values(1,1);  
insert into TM values(2,2);  
insert into TM values(3,3);
```

```
tm  
— delta_000001_000001_0000  
  └ 00000_0  
— delta_000002_000002_0000  
  └ 00000_0  
— delta_000003_000003_0000  
  └ 00000_0
```

Transaction Manager

- Transaction State
 - Open, Committed, Aborted
- Reader at Snapshot Isolation
 - A snapshot is the state of all transactions
 - High Water Mark + List of Exceptions

```
tm
  — delta_0000001_0000001_0000
    \— 000000_0
  — delta_0000002_0000002_0000
    \— 000000_0
  — delta_0000003_0000003_0000
    \— 000000_0
```

- ◆ Atomicity & Isolation

Full CRUD

- No in-place Delete - Append-only file system
- Isolate readers from writers

ROW_ID

- CREATE TABLE acidtbl (a INT, b STRING) STORED AS ORC TBLPROPERTIES ('transactional'='true');

Metadata Columns	original_write_id bucket_id row_id current_write_id	ROW_ID
User Columns	col_1: a : INT col_2: b : STRING	

Create

- `INSERT INTO acidtbl (a,b) VALUES (100, "foo"), (200, "xyz"), (300, "bee");`

ROW_ID	a	b
{ 1, 0, 0 }	100	"foo"
{ 1, 0, 1 }	200	"xyz"
{ 1, 0, 2 }	300	"bee"

`delta_00001_00001/bucket_0000`

Delete

- DELETE FROM acidTbl where a = 200;

ROW_ID	a	b
{ 1, 0, 0 }	100	"foo"
{ 1, 0, 1 }	200	"xyz"
{ 1, 0, 2 }	300	"bee"

ROW_ID	a	b
{ 1, 0, 1 }	null	null

delete_delta_00002_00002/bucket_0000

delta_00001_00001/bucket_0000

- ◆ Readers skip deleted rows

Update

- Update = delete + insert
- ◆ UPDATE acidTbl SET b = “bar” where a = 300;

ACID_PK	a	b
{ 1, 0, 0 }	100	“foo”
{ 1, 0, 1 }	200	“xyz”
{ 1, 0, 2 }	300	“bee”

delta_00001_00001/bucket_0000

ACID_PK	a	b
{ 2, 0, 0 }	300	“bar”

delta_00003_00003/bucket_0000

ACID_PK	a	b
{ 1, 0, 2 }	null	null

delete_delta_00003_00003/bucket_0000

Read

- Ask Transaction Manager for Snapshot Information
 - Decide which deltas are relevant
- Take all the files in delta_x_x/ and split them into chunks for each processing Task to work with
- Localize all delete events from each delete_deleta_x_x/ to each task
 - Highly Compressed with ORC
- Filter out all Insert events that have matching delete events
 - Requires an Acid aware reader – thus AcidInputFormat

Design - Compactor

- More Update operations = more delete events – make reads more expensive
- Insert operations don't add read overhead

Design - Compactor

- Compactor rewrites the table in the background
 - Minor compaction - merges delta files into fewer deltas
 - Major compactor merges deltas with base - more expensive
 - This amortizes the cost of updates and self tunes the tables
 - Makes ORC more efficient - larger stripes, better compression
- Compaction can be triggered automatically or on demand
 - There are various configuration options to control when the process kicks in.
 - Compaction itself is a Map-Reduce job
- Key design principle is that compactor does not affect readers/writers
- Cleaner process – removes obsolete files
- Requires Standalone metastore

Merge Statement – SQL Standard 2011 (Hive 2.2)

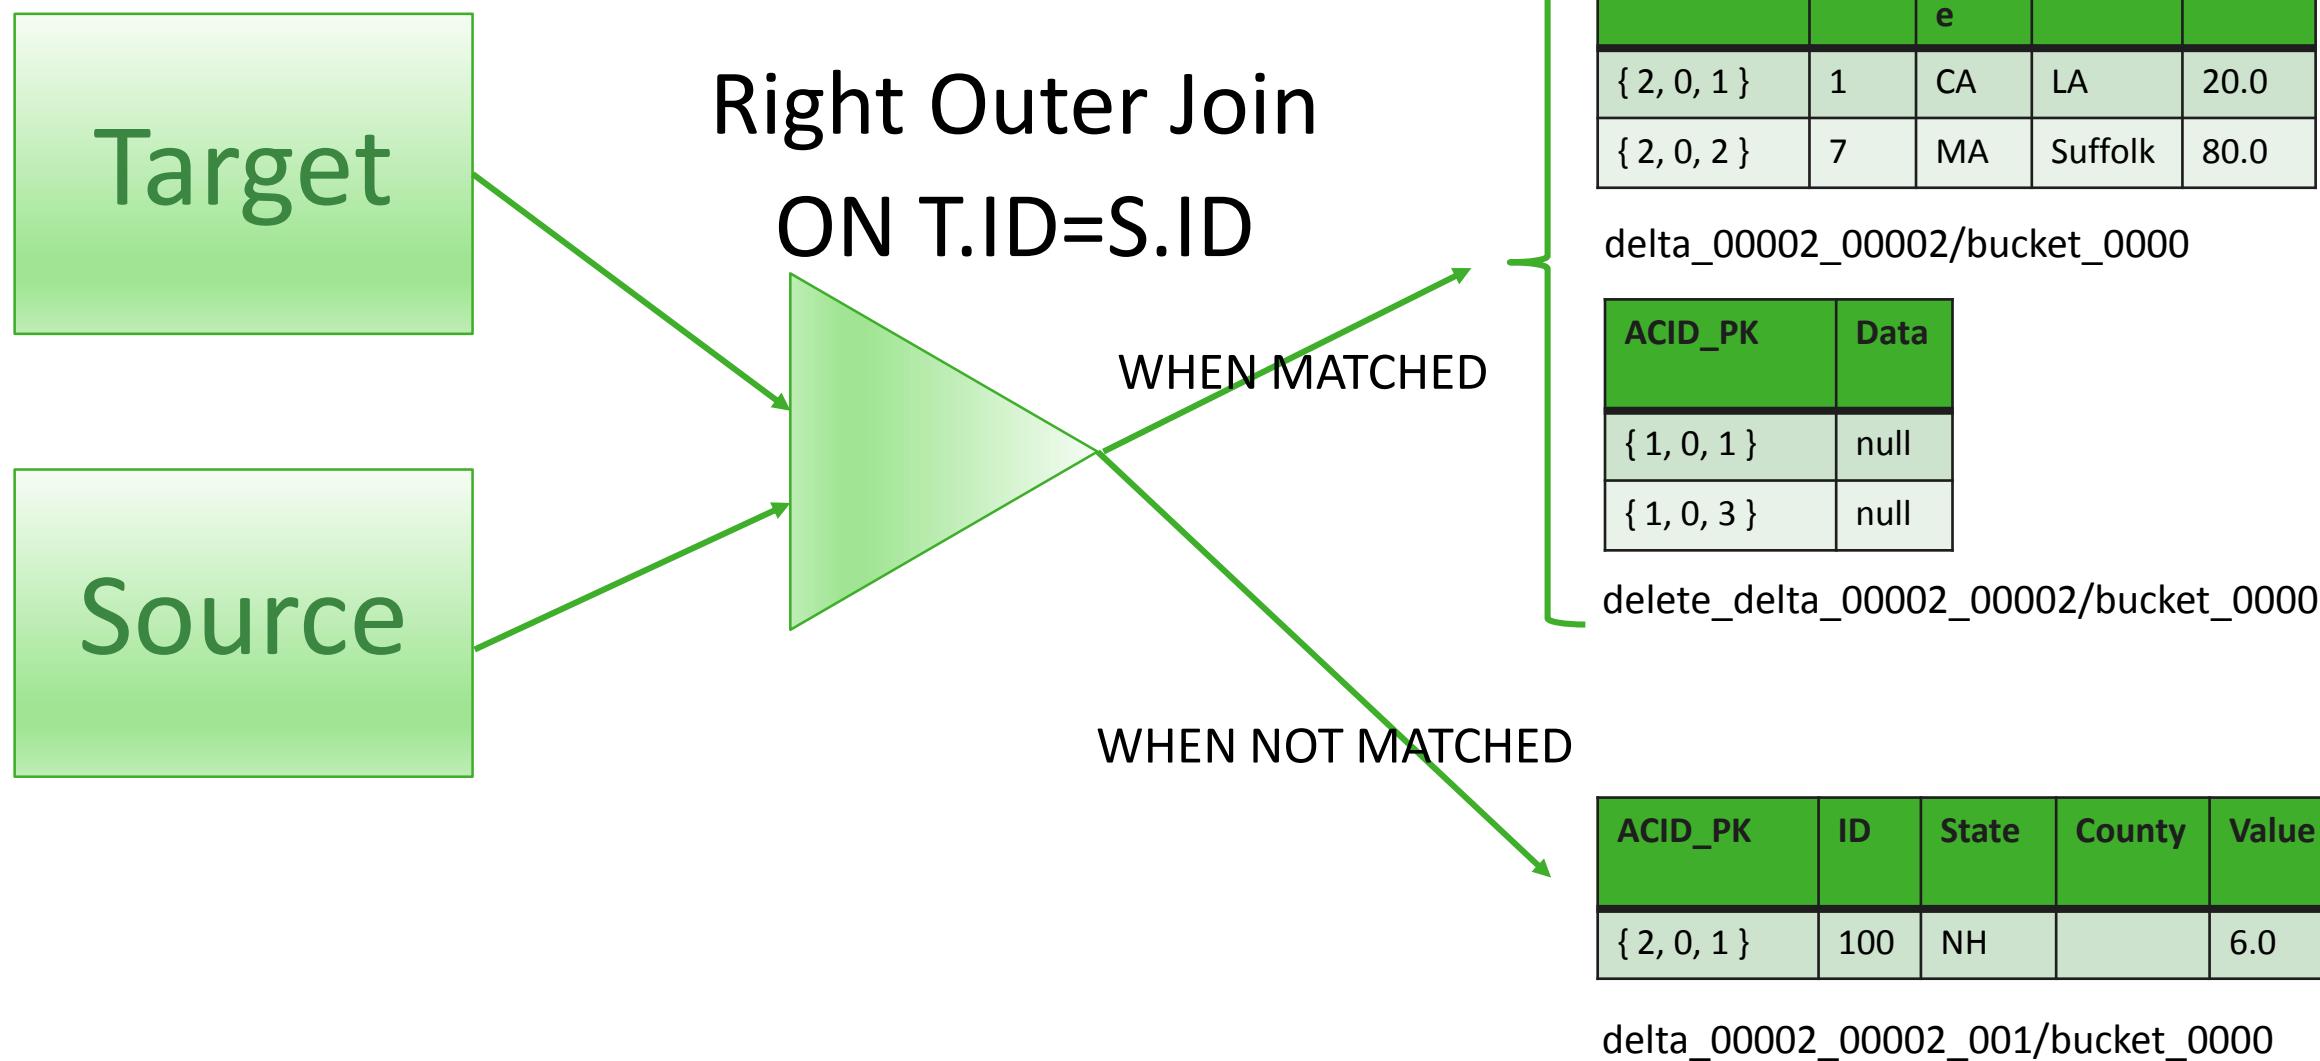
ID	State	County	Value
1	CA	LA	19.0
2	MA	Norfolk	15.0
7	MA	Suffolk	50.15
16	CA	Orange	9.1

ID	State	Value
1		20.0
7		80.0
100	NH	6.0

```
MERGE INTO TARGET T
  USING SOURCE S ON T.ID=S.ID
  WHEN MATCHED THEN
    UPDATE SET T.Value=S.Value
  WHEN NOT MATCHED
    INSERT (ID,State,Value)
    VALUES(S.ID, S.State, S.Value)
```

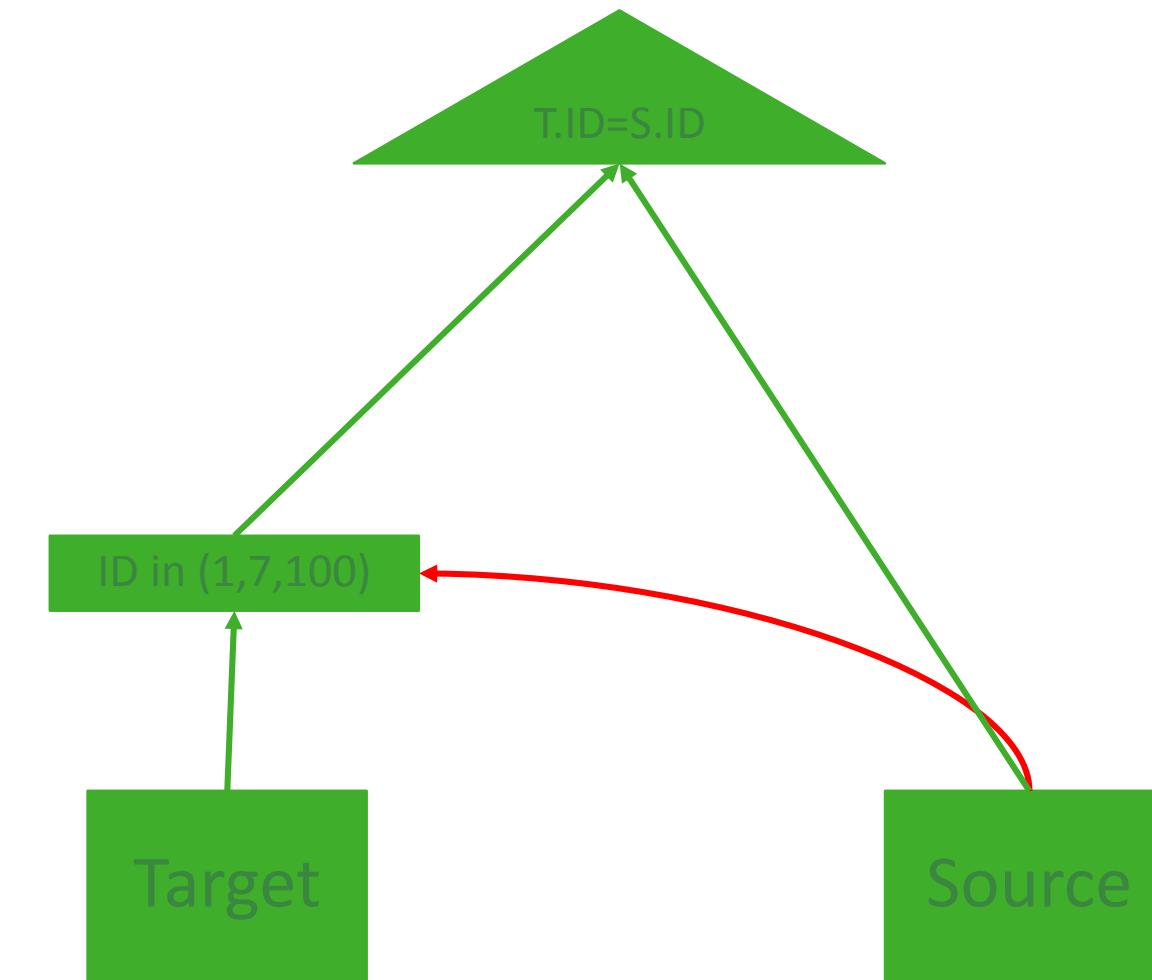
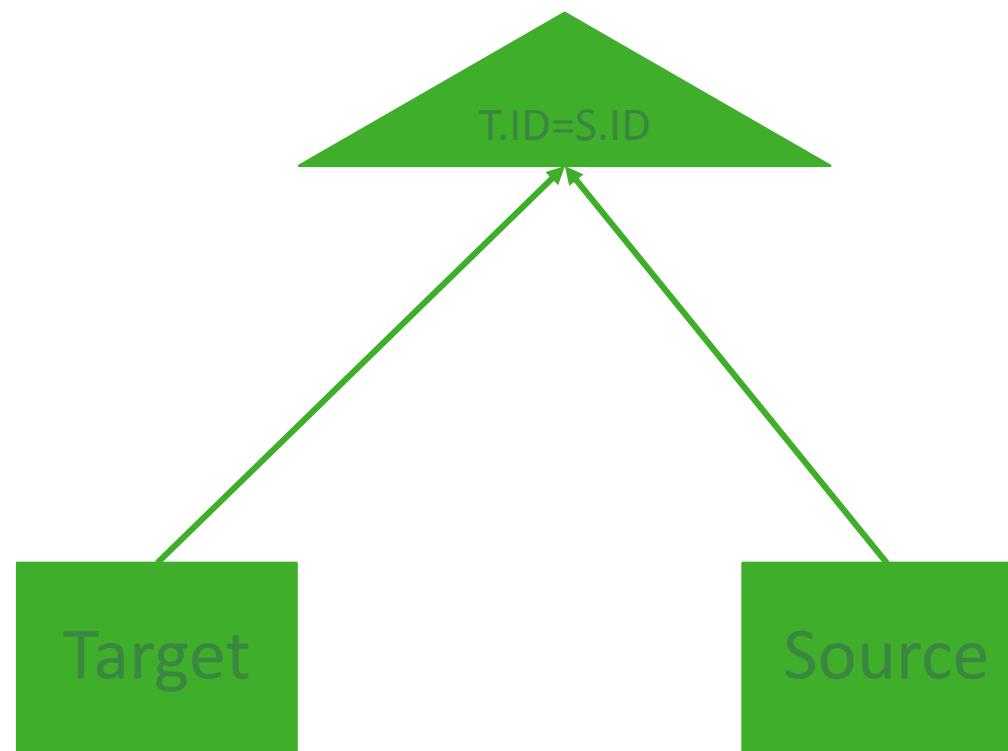
ID	State	County	Value
1	CA	LA	20.0
2	MA	Norfolk	15.0
7	MA	Suffolk	80.0
16	CA	Orange	9.1
100	NH	null	6.0

SQL Merge



Merge Statement Optimizations

- Semi Join Reduction
 - aka Dynamic Runtime Filtering
 - On Tez only



Design - Concurrency

- Inserts are never in conflict since Hive does not enforce unique constraints
- Write Set tracking to prevent Write-Write conflicts in concurrent transactions
- Lock Manager
 - DDL operations acquire eXclusive locks – metadata operations
 - Read operations acquire Shared locks

Tooling

- SHOW COMPACTIONS
 - Hadoop Job ID
- SHOW TRANSACTIONS
- SHOW LOCKS
 - What a lock is blocked on
- ABORT TRANSACTIONS txnid1, txnid2....

Other Subsystems

- Result Set Caching
 - Is it valid for current reader?
- Materialized Views
 - Incremental View Maintenance
- Spark
 - HiveWarehouseConnector: HS2 + LLAP

Streaming Ingest API

- Connection – Hive Table
 - Begin transaction
 - Commit/Abort transaction
 - org.apache.hive.streaming.StreamingConnection
- Writer
 - Write records
 - org.apache.hive.streaming.RecordWriter
- Append Only via this API
 - Update/Delete via SQL
- Optimized for Write operations
- Requires more aggressive Compaction for efficient reads
- Supports dynamic partitioning in a single transaction

Limitations

- Transaction Manager
 - State is persisted in the metastore RDBMS
- Begin/Commit/Abort
 - Metastore calls

Future

Future Work

- Multi statement transactions, i.e. BEGIN TRANSACTION/COMMIT/ROLLBACK
- Performance
 - Smarter Compaction
- Finer grained concurrency management/conflict detection
- Read Committed w/Lock Based scheduling
- Better Monitoring/Alerting
- User define Primary Key
 - Transactional Tables sorted on PK

Further Reading

Etc

- Documentation
 - <https://cwiki.apache.org/confluence/display/Hive/Hive+Transactions>
 - <https://cwiki.apache.org/confluence/display/Hive/Streaming+Data+Ingest+V2>
- Follow/Contribute
 - <https://issues.apache.org/jira/browse/HIVE-14004?jql=project%20%3D%20HIVE%20AND%20component%20%3D%20Transactions>
- user@hive.apache.org
- dev@hive.apache.org

Credits

- Alan Gates
- Sankar Hariappan
- Prasanth Jayachandran
- Eugene Koifman
- Owen O'Malley
- Saket Saurabh
- Sergey Shelukhin
- Gopal Vijayaraghavan
- Wei Zheng



Thank You